



**Snowflake**

## **Business Scenarios**

Components

Know Snowflake

SnowSQL

Loading Data

3rd Party Connections

## **Agenda.**

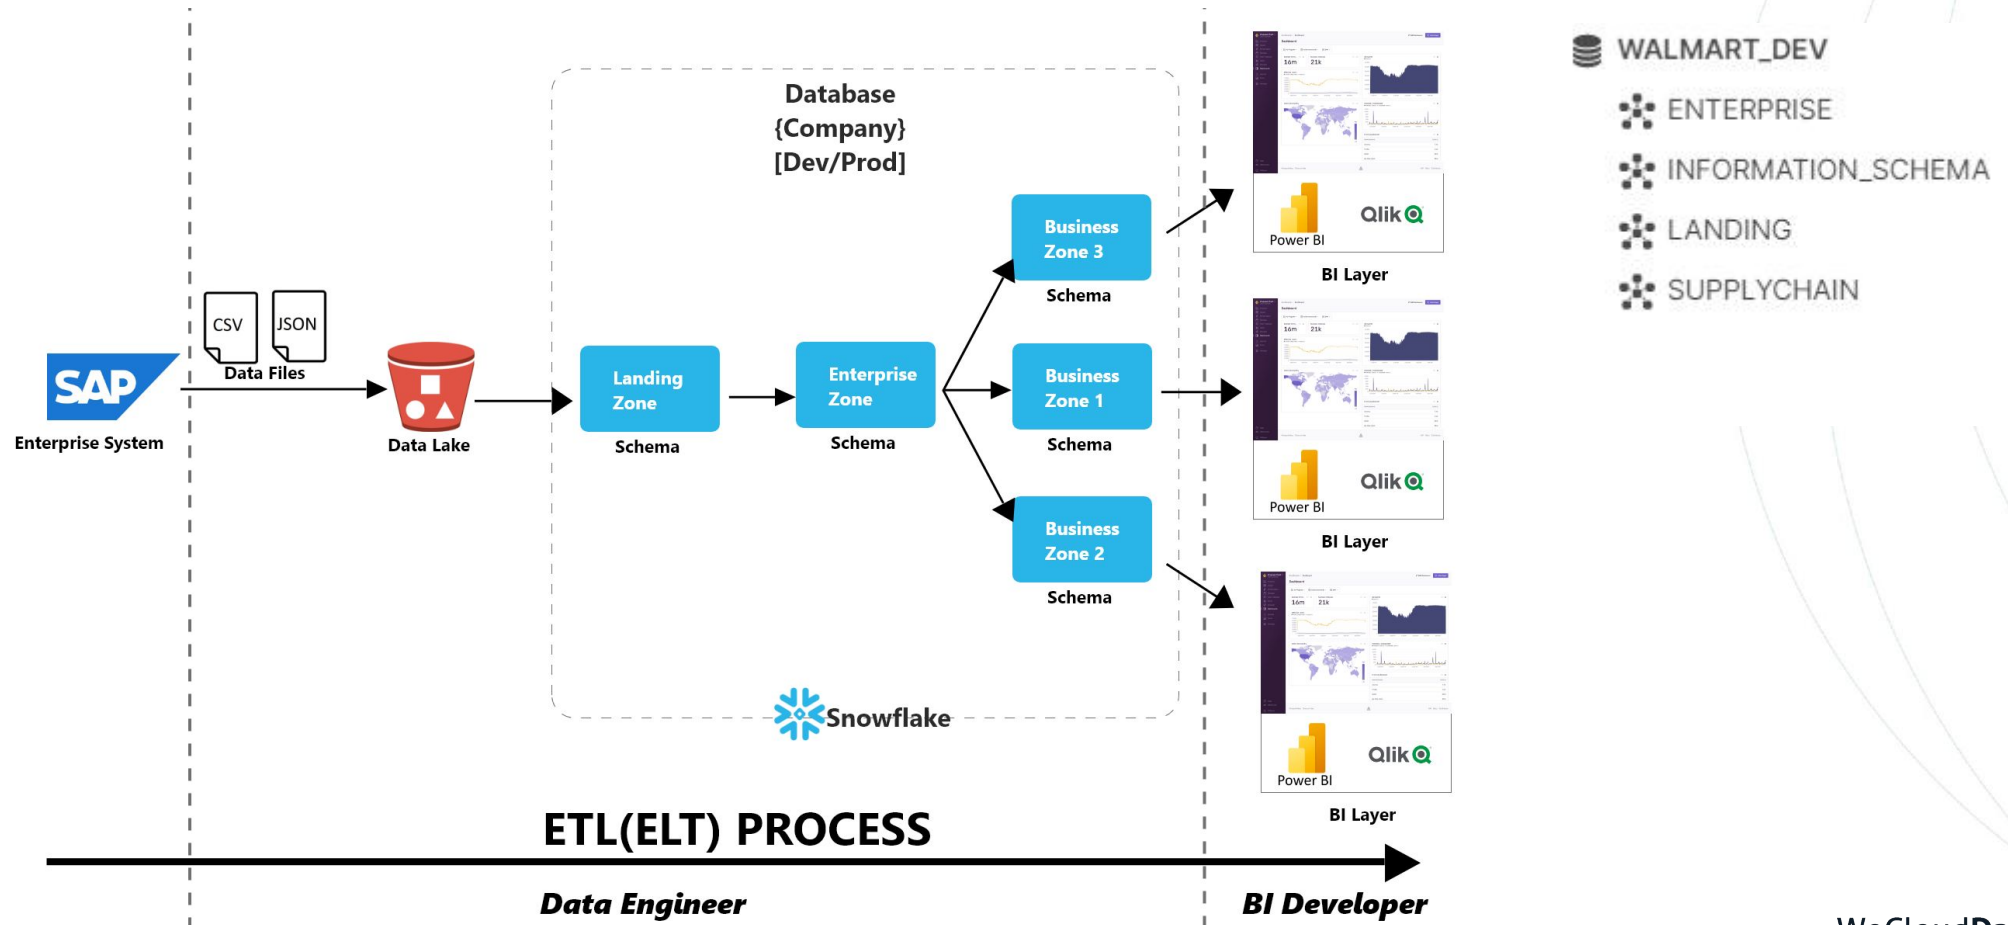


# Snowflake Usage

## Business Scenarios

### Dataflow

Snowflake used as a Data warehouse in a BI platform.



Business Scenarios

## **Components**

Know Snowflake

SnowSQL CLI

Loading Data

3rd Party Connections

## **Agenda.**



# Components in Snowflake

## Components

- **Database**
  - Database is an organization's data sets. Dev and Prod use different databases.
- **Schema**
  - schema shows a logical view of an database. It constructs views and tables. It represent a set of data in certain purpose.
- **Table and views**
  - Contain structured data
- **Worksheet**
  - Editor to write and run SQL.
- **Warehouse**
  - Engine to run queries.
- **Roles**
  - Access and operation permission group.
- **History**
  - Historical queries.



# Warehouse

## Components

### ❖ Create Warehouse in Console

Warehouses								+ Warehouse
2 Warehouses		Q Search Name		Status All	Size All	C		
NAME ↑	STATUS	SIZE	TYPE	RUNNING	QUEUED	OWNER	CREATED	
DEMO	Suspended	X-Small	Standard	0	0	ACCOUNTAD...	5 months ago	...

### ❖ Create Warehouse in Query

```
CREATE [ OR REPLACE ] WAREHOUSE [ IF NOT EXISTS ] <name>
WITH WAREHOUSE_SIZE = XSMALL | SMALL | MEDIUM | LARGE | XLARGE | XXLARGE | XXXLARGE |
X4LARGE | X5LARGE | X6LARGE
MAX_CLUSTER_COUNT = <num>
MIN_CLUSTER_COUNT = <num>
... .
;

USE WAREHOUSE <name>;
```

Useful Link: <https://docs.snowflake.com/en/sql-reference/sql/create-warehouse.html>




Try in system



# Database

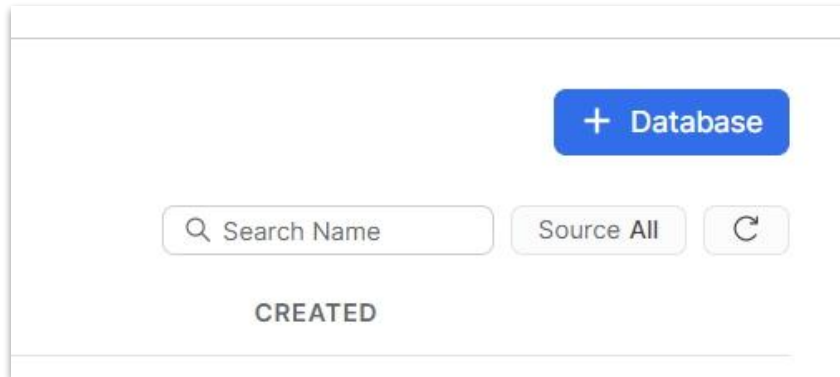
## Components

**Definition:** Database is an organization's data sets. Dev and Prod use different databases.



## Creation

- Console
  - Creation



- Query

```
CREATE DATABASE IF NOT EXISTS DEMO_DB;  
  
USE DATABASE DEMO_DB;
```




Try in system



# Schema Components

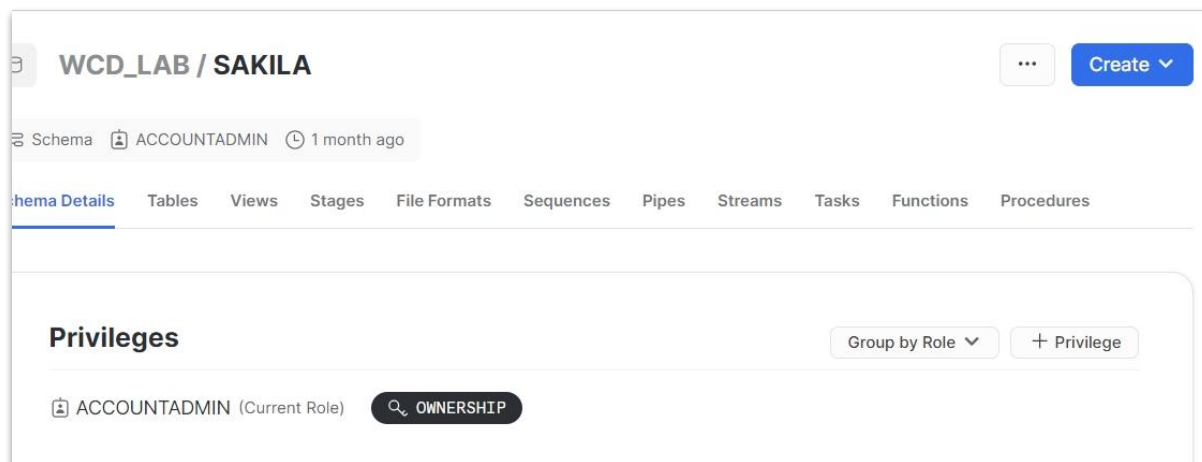
**Definition:** Where you split your database logically.



## Creation

- Console
  - Creation

- Query



```
CREATE SCHEMA IF NOT EXISTS DEMO_DB.SCHMA;
```



Try in system





# Schema

## Components

### Data Transit between schemas:



- **Create**

```
CREATE TABLE IF NOT EXISTS enterprise.table  
AS (SELECT col_1, col_2 FROM landing.table);
```

- **Insert**

```
INSERT INTO enterprise.table (col_1, col_2)  
(SELECT col_1, col_2 FROM landing.table);
```



Try in system



# Tables and Views

## Components

Name	Definition	Usages
<b>Regular Table</b>	Regular Table is a collection that contains the data physically, which means the data take store space in hardwares.	Contain the formal tables in the database.
<b>Transient Table</b>	Transient tables are similar to regular tables with the main difference that they do not have a Fail-safe period.	It is mainly used as a staging table in ETL process.
<b>Temporary Table</b>	Temporary table is a physical table, but it only exist within the <b>session</b> in which they were created and persist only for the remainder of the session.They are not visible to other users or sessions.	Temp table is rarely used on ETL process. It is only used when you want to create some simple and temporary staging table in ETL; or when you need to do some testing but don't want impact the entire database. <a href="#">[Try in system]</a>
<b>View</b>	A view allows the result of a query to be accessed as if it were a table.	It is used when you only need the query result, but don't need it contains data permanently. Mostly in ETL process.
<b>Materialized View</b>	A materialized view is a pre-computed data set derived from a query and stored for later use. Querying a materialized view is faster than a view.	It is expensive than a regular view, so it is used when: <ul style="list-style-type: none"><li>• you need faster and significant processing;</li><li>• the result is small size data;</li><li>• The aggregates that take a long time to calculate.</li></ul>



# Regular , Transient, Temporary Table

## Components

- **Table Information**

- ❖ **Show all the tables in a schema**

```
SHOW TABLES IN <database>.<schema>;
```

- ❖ **Display a specific table information**

```
DESCRIBE TABLE <database>.<schema>.<table>;
```

- ❖ **Show the columns in a table**

```
SHOW COLUMNS IN <database>.<schema><table>;
```





# Regular , Transient, Temporary Table

## Components

- **Table Creation**

- ❖ **Create an table**

```
CREATE TABLE [TRANSIENT TABLE, TEMP TABLE].....
```

- ❖ **Create an empty table**

```
CREATE [OR REPLACE] TABLE [IF NOT EXISTS] <table name>
(<col1_name> <col_type>,
 <col2_name> <col_type>,
 .....
 <coln_name> <col_type>)
[CLUSTER BY (col1_name,col2_name)];
```

```
CREATE TABLE IF NOT EXISTS enterprise.orders
(col_1 int,
col_2 int)
CLUSTER BY (col_1,col_2);
```

- ❖ **Create an empty table by copying another table structure**

```
CREATE [OR REPLACE] TABLE [IF NOT EXISTS] <table name>
LIKE <another table name>
[CLUSTER BY (col1_name,col2_name)];
```

```
CREATE TABLE IF NOT EXISTS enterprise.orders_new
LIKE enterprise.orders
CLUSTER BY (col_1);
```





# Regular , Transient, Temporary Table

## Components

- **Table Creation**

- ❖ **Create a table by clone another table totally**

```
CREATE [OR REPLACE] TABLE [IF NOT EXISTS] <table name>  
CLONE <another table>  
[CLUSTER BY (col1_name,col2_name)];
```

```
CREATE OR REPLACE TABLE enterprise.orders_new  
CLONE enterprise.orders;
```

- ❖ **Create an table by conditionally copying another table**

```
CREATE [OR REPLACE] TABLE [IF NOT EXISTS] <table name>  
(<col1_name> <col_type>,  
 <col2_name> <col_type>,  
 <col3_name> <col_type>)  
AS  
  (SELECT col1,col2,col1 FROM <another table>)  
[CLUSTER BY (col1_name,col2_name)];
```

```
CREATE TABLE IF NOT EXISTS enterprise.city_name  
(CITY_NAME VARCHAR)  
AS (SELECT CTY_NAME FROM enterprise.citys);
```









# Regular , Transient, Temporary Table

## Components

- **Table Creation**

 **WCD\_LAB / SAKILA** ... Create ▾

 Schema  ACCOUNTADMIN  1 month ago

Schema Details **Tables** Views Stages File Formats Sequences Pipes Streams Tasks Functions Procedures

**16 Tables**  All Tables ↻

NAME ↑	TYPE	OWNER	ROWS	BYTES	CREATED
--------	------	-------	------	-------	---------

Useful Link: <https://docs.snowflake.com/en/sql-reference/ddl-table.html>



Try in system



# Table

## Components

- **Table Modification**

- ❖ **Change name**

- ```
ALTER TABLE [IF EXISTS] <table name> RENAME TO <new table name>;
```

- ❖ **Add and change Column**

- ```
ALTER TABLE [IF EXISTS] <table name> ADD COLUMN <col_name> <col_type>;
```

- ```
ALTER TABLE [IF EXISTS] <table name> RENAME COLUMN <col_name> to <new col_name>;
```

- ```
ALTER TABLE [IF EXISTS] <table name> ALTER COLUMN <col_name> SET TYPE <new data_type>;
```

- **Drop Table**

- ```
DROP TABLE [IF EXISTS] <table name> ;
```

- ```
UNDROP TABLE [IF EXISTS] <table name> ;
```





# View and Materialized view

## Components

- **Table Information**

- ❖ **Show all the views in a schema**

```
SHOW VIEWS IN <database>.<schema>;
```

- ❖ **Display a specific view information**

```
DESCRIBE VIEW <database>.<schema>.<view>;
```

- ❖ **Show the columns in a view**

```
SHOW COLUMNS IN <database>.<schema><view>;
```







# View and Materialized Views

## Components

- **View Creation**

- ❖ **Create an view by SELECT**

```
CREATE [OR REPLACE] [MATERIALIZED] VIEW [IF NOT EXISTS] <view name>
(<col1_name> <col_type>,
 <col2_name> <col_type>,
 <col3_name> <col_type>)
AS
(SELECT col1,col2,col1 FROM <another table>)
[CLUSTER BY (col1_name,col2_name)];
```

- ❖ **Drop View**

```
DROP [MATERIALIZED] VIEW [IF EXISTS] <view name> ;
```





# View and Materialized Views

## Components

### Deciding When to Create a Materialized View or a Regular View🔗

In general, when deciding whether to create a materialized view or a regular view, use the following criteria:

- Create a materialized view when **all** of the following are true:
  - The query results from the view don't change often. This almost always means that the underlying/base table for the view doesn't change often, or at least that the subset of base table rows used in the materialized view don't change often.
  - The results of the view are used often (typically significantly more often than the query results change).
  - The query consumes a lot of resources. Typically, this means that the query consumes a lot of processing time or credits, but it could also mean that the query consumes a lot of storage space for intermediate results.
- Create a regular view when **any** of the following are true:
  - The results of the view change often.
  - The results are not used often (relative to the rate at which the results change).
  - The query is not resource intensive so it is not costly to re-run it.



# Tables and Views Comparison

## Components

Name	Speed	Cost
Cache	Very Fast	Very Expensive
Materialized View	Fast	Expensive
Regular Table	OK	OK
Regular View	Slow	Low



Try in system



# Tables – Data type

## Components

Name	Definition	Example
<b>INT</b>	Integer	product_key, Is_flag (converted), qty (sometimes)
<b>NUMERIC(38,2)</b>	A number, system decide 0 ~ 2 digitals	sales, avg_xxx, sum_xxx
<b>VARCHAR (XXX)</b>	Characters: If you no number (xxx) is defined, system will adjust automatically; If defined, the length of character will be fixed. Usually don't define.	strings
<b>BOOLEAN</b>	Yes or No.	Is_flag
<b>DATE</b>	Date format: 2021-01-01	date
<b>TIMESTAMP</b>	Date and time: 2022-06-30 10:47:18.480 -0700	table update time



Try in system



# Worksheet

## Components

WCD create s3 stage ▾

Queries folder

Worksheets Databases

Pinned (0)

No pinned objects

Q All Objects ...

SNOWFLAKE

databases

WALMART\_DEV

WCD\_LAB

WCD\_LCT4

WCD\_LECTURE

WCD\_SQL

query editor

WALMART\_DEV.ENTERPRISE ▾

23 FIELD\_DELIMITER = ',';

24

25

26

27 GRANT CREATE STAGE ON SCHEMA ENTERPRISE to ROLE accountadmin;

28 GRANT USAGE ON INTEGRATION S3\_INT\_WCD\_LCT1 to ROLE accountadmin;

29

30 CREATE OR REPLACE STAGE WCD\_LCT1\_STAGE

31 STORAGE\_INTEGRATION = S3\_INT\_WCD\_LCT1

32 URL='s3://snowflake-stage-bucket-eric'

33 FILE\_FORMAT = CSV\_COMMA;

34

35

36 LIST @WCD\_LCT1\_STAGE;

37

38 COPY INTO walmart\_dev.enterprise.city from @WCD\_LCT1\_STAGE/city.csv FILE\_FORMAT =(SKIP\_HEADER=1);

ACCOUNTADMIN COMPUTE\_WH Share

role warehouse run

Updated 1 week ago

Objects Editor Results Chart

file status rows\_parsed rows\_loaded error\_limit errors\_seen

s3://snowflake-stage-bucket-eric/city.csv LOADED 319 319 1 0

Query Details ...

Query duration 1.9s



Try in system



# Role

## Components

### ❖ The purpose of the Role

A Role can decide an account can use which:

- Database
- Schema
- Table(rarely happen)
- Warehouse

### ❖ Create a Role

```
CREATE [ OR REPLACE ] ROLE [ IF NOT EXISTS ] <role_name>;
```

```
GRANT USAGE DATABASE <database_name> TO ROLE <role_name>;
```



Try in system

Business Scenarios

Components

**Know Snowflake**

SnowSQL CLI

Loading Data

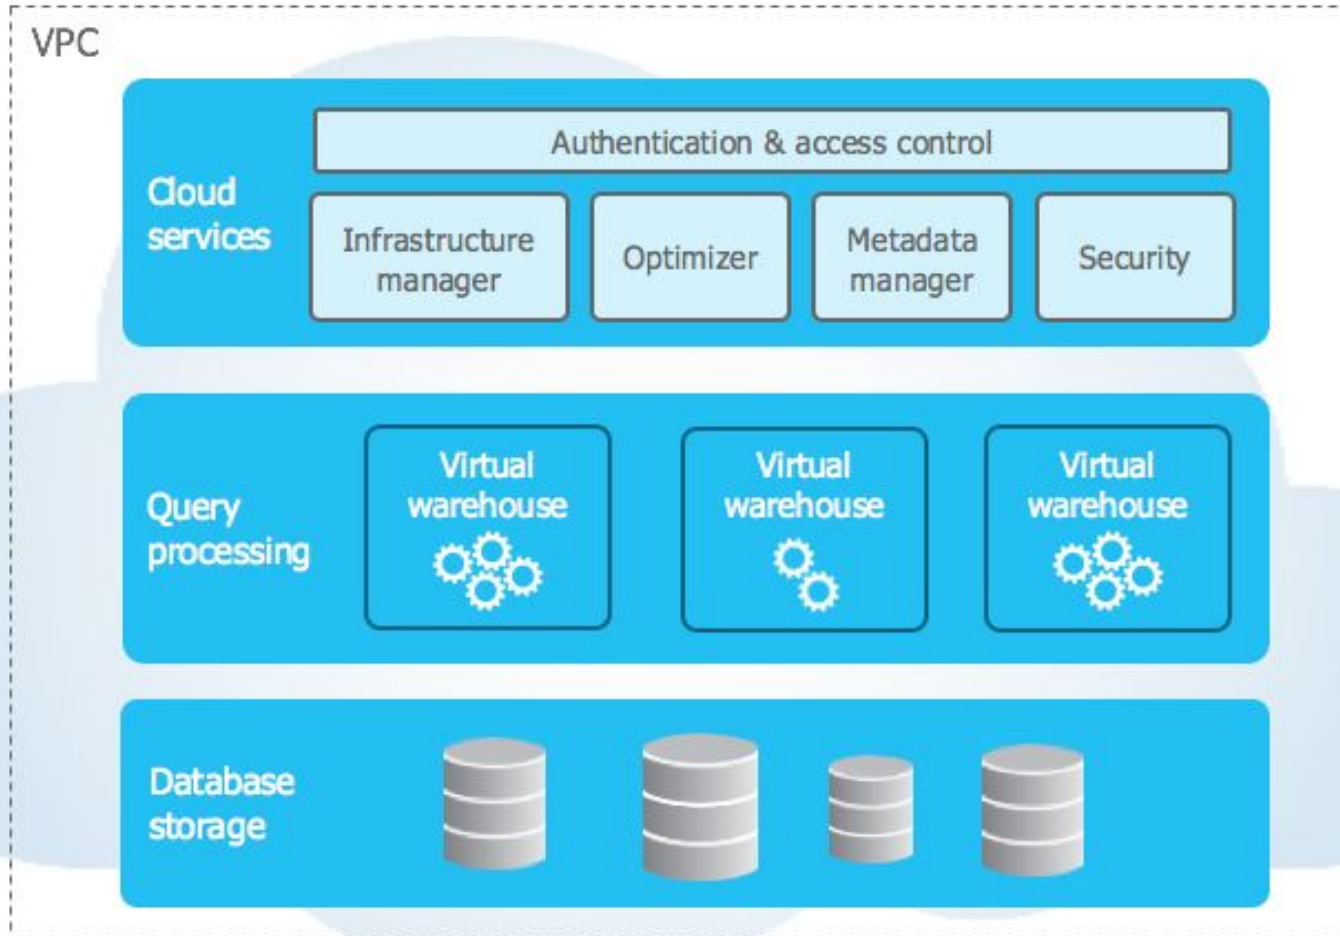
3rd Party Connections

**Agenda.**



# Snowflake Architecture

Know Snowflake







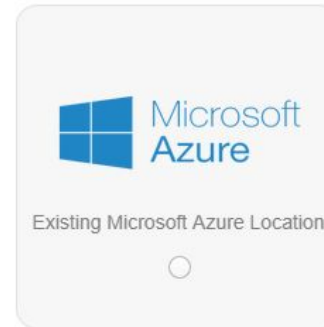
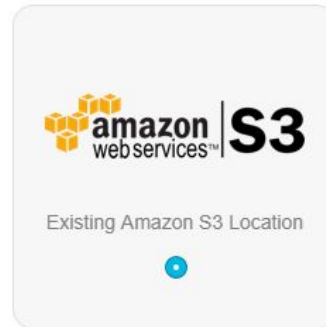
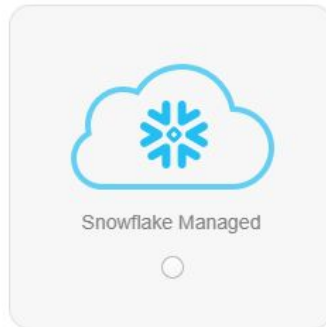
# Snowflake Storage

Know Snowflake

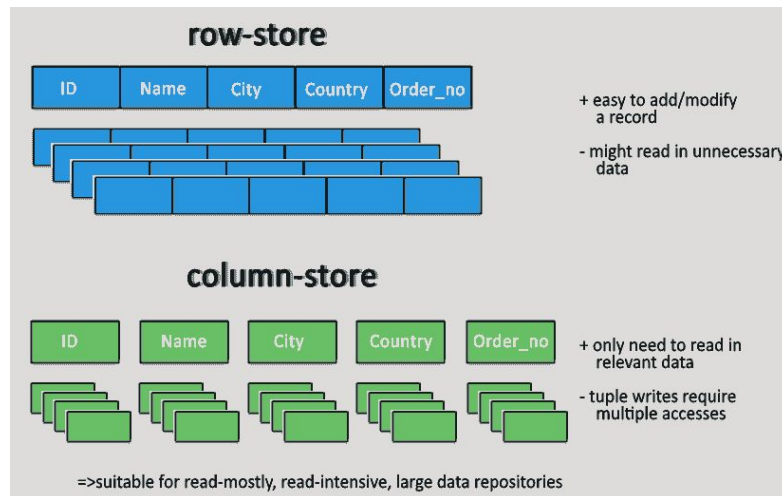
## ❖ The freedom to store your data

### Create Stage

Choose a location for files to be staged



## ❖ The Data store in column



- Good at where clause with keys
- Good at joining
- Primary key, foreign key will not be useful in snowflake
- updating will be expensive
  - So avoid 'update'



Try in system



# Snowflake Pricing

Know Snowflake

- **Storage**
  - \$23 USD per compressed TB each month of data stored in US.
- **Computing**
  - Compute costs are \$0.00056 **per second** for **each credit** consumed on Snowflake Standard Edition.

Snowflake Warehouse Sizes and Credit Usage per Hour										
Size	X-Small	Small	Medium	Large	X-Large	2X-Large	3X-Large	4X-Large	5X-Large (In preview)	6X-Large (In preview)
Credit Usage per Hour	1	2	4	8	16	32	64	128	256	512

**Example:** *X-Large*

$16 \times 3600 \times 0.00056 = 32.256$  dollars

Business Scenarios

Components

Know Snowflake

**SnowSQL CLI**

Loading Data

3rd Party Connections

**Agenda.**



# installation on Linux (For Ubuntu)

## SnowSQL CLI

### 1. Go to Download Folder

*cd ~/Downloads*

### 1. Download the SnowSQL CLI Application

*wget*

*[https://sfc-repo.snowflakecomputing.com/snowsql/bootstrap/1.2/linux\\_x86\\_64/snowflake-snowsql-1.2.21-1.x86\\_64.rpm](https://sfc-repo.snowflakecomputing.com/snowsql/bootstrap/1.2/linux_x86_64/snowflake-snowsql-1.2.21-1.x86_64.rpm)*

### 1. Install RPM Package Manager

*sudo apt-get update*

*sudo apt-get install alien -y*

### 1. Install the SnowSQL

*sudo alien -i snowflake-snowsql-1.2.21-1.x86\_64.rpm*

### 1. Check version

*snowsql -v*

Useful Link: <https://docs.snowflake.com/en/user-guide/snowsql-install-config.html#installing-snowsql-on-linux-using-the-rpm-package>



# installation on Mac

## SnowSQL CLI

### Installing SnowSQL on macOS Using Homebrew Cask%

Homebrew Cask<sup>®</sup> is a popular extension of Homebrew<sup>®</sup> used for package distribution, installation, and maintenance. There is no separate SnowSQL installer to download. If Homebrew Cask is installed on your macOS platform, you can install Snowflake directly.

Run the `brew install` command, specifying `snowflake-snowsql` as the cask to install:

```
$ brew install --cask snowflake-snowsql
```

### Configuring the Z Shell Alias (macOS Only)%

If Z shell (also known as zsh) is your default terminal shell, set an alias to the SnowSQL executable so that you can run SnowSQL on the command line in Terminal. The SnowSQL installer installs the executable in `/Applications/SnowSQL.app/Contents/MacOS/snowsql` and appends this path to the PATH or alias entry in `~/.profile`. Because zsh does not normally read this file, add an alias to this path in `~/.zshrc`, which zsh **does** read.

To add an alias to the SnowSQL executable:

1. Open (or create, if missing) the `~/.zshrc` file.
2. Add the following line:

```
alias snowsql=/Applications/SnowSQL.app/Contents/MacOS/snowsql
```

3. Save the file.



[Back to top](#)

Useful Link: <https://docs.snowflake.com/en/user-guide/snowsql-install-config.html#installing-snowsql-on-linux-using-the-rpm-package>



# Connect SnowSQL with Snowflake

## SnowSQL CLI

### 1. Use command with account-name and username

```
snowsql -a <account-name> -u <username>
```

### 2. Input Password

```
Failed to initialize log. No logging is enabled: [Errno 13] Permission denied: '/home/snowsql_rt.log'  
Password: 
```

### 3. SnowSQL start running

```
* SnowSQL * v1.2.21  
Type SQL statements or !help  
WCDDE3#COMPUTE_WH@(no database).(no schema)>
```



# Connect SnowSQL with Snowflake

## SnowSQL CLI

https://ozb44782.us-east-1.snowflakecomputing.com/console/login#/  
Account Name

User Name

Sign in to Snowflake

Username  
snowflake050701

Password  
.....

[Forgot password](#)

Sign in



# SnowSQL config

SnowSQL CLI

## Why Config?

- Connection

→ Default Connection Setting:

```
#If a connection doesn't specify a value, it will default to these
#
#accountname = defaultaccount
#region = defaultregion
#username = defaultuser
#password = defaultpassword
#dbname = defaultdbname
#schemaname = defaultschema
#warehousename = defaultwarehouse
#rolename = defaultrolename
#proxy_host = defaultproxyhost
#proxy_port = defaultproxyport
```

→ Specific Connection Setting:

```
[connections.wcd]
#Can be used in SnowSql as #connect example

accountname = jo56804.ca-central-1.aws
username = WCDDE3
password = xxxxxxxxxxxxxxxxx
```

## Config file

Location: `~/.snowsql/config`

Open: `nano ~/.snowsql/config`

## Connection:

default setting: `snowsql .....`

Specific setting: `snowsql -c wcd`



Try in system





# SnowSQL Command

SnowSQL CLI

→ Options:

## SnowSQL command

**snowsql** **-c** wcd **-o** variable\_substitution=true **-D** q\_table=enterprise.city **-D** cty\_id=100001 **-f** script.sql

```
Use database walmart_Dev;  
  
select * from &{q_table}  
where cty_id = &{cty_id};  
query in script.sql
```

**-c** Connection

**-q** Query

**-o** Option. Here must HAVE `variable_substitution=true` , otherwise, **-D** will not work.

**-D** Variables in script file

**-f** File of script



Try in system

Business Scenarios

Components

Know Snowflake

SnowSQL CLI

**Loading Data**

3rd Party Connections

**Agenda.**



# Load Data with Query – Steps

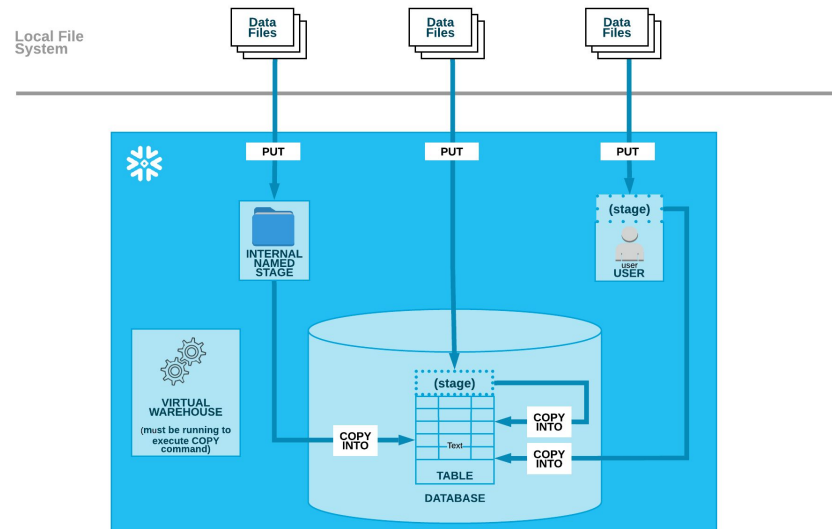
Loading data

## Step 1

Upload (i.e. **stage**) one or more data files to a Snowflake stage (named internal stage or table/user stage) using the **PUT** command.

## Step 2

Use the **COPY INTO** <table> command to load the contents of the staged file(s) into a Snowflake database table.



Useful Link: <https://docs.snowflake.com/en/user-guide/data-load-local-file-system.html>



# Staging Area

Loading data

Before a file being loaded into a table, it must be loaded into a Staging Area first. There are 2 important types of Staging Areas — User Stages and Named Stages:

## User Stages

By default, each user and table in Snowflake is automatically allocated an internal stage for staging data files to be loaded. To list the file under the user stage use '*LIST @~;*'

## Named Stages

Named stages are database objects that provide the greatest degree of flexibility for data loading. Named stages are optional but recommended when you plan regular data loads that could involve multiple users and/or tables. Usually it is stages integrate with other data lake, such as S3.



# Load Data From User Stages

Loading data

1. **Load data into Staging Area** (*The process must be done in CLI*)

```
snowsql -c wcd -q "PUT file://countries.csv @~"
```

-----  
@~: is the User stage.

file://cities.csv is the file dir in the local system.

1. **Copy Data from Staging area to the table** (*The process works in both Editor or CLI*)

```
COPY INTO schma.countries FROM @~/countries.csv [FILE_FORMAT = (TYPE = CSV FIELD_DELIMITER = ','  
SKIP_HEADER = 1)];
```

```
COPY INTO schma.countries FROM (SELECT $1, $2 FROM @~/cities.csv);
```

@~: is the User stage

/cities.csv is the file dir in the stage.



# Load Data From Named Stages (S3 Integration)

Loading data

A stage integrated with S3 can upload or download a file from S3 directly. In order to make this, you need to set:

- IAM role in AWS which allows snowflake access to S3.
- Create a INTEGRATION at Snowflake.
- Create a STAGE with such INTEGRATION.
- The detail steps to create a stage can be found from this [Introduction](#).

The lecture scripts can be found [here](#).



## Demo for local stage

Create a Database WALMART\_DEV, a schema ENTERPRISE, a table CITY:

**1. Load data to local stage “@~” from file with “PUT”, and “COPY” the stage file into the city table.**

- a. Create database, schema, and table with [this query](#);
- b. download CLI on your Linux server, and config the connection;
- c. Load city data([download from here](#)) in local stage “@~”, and then copy from “@~” to city table with command:

```
snowsql -c wcd -q "put file://city.csv @~; copy into walmart_dev.enterprise.city from @~/city.csv.gz FILE_FORMAT =(SKIP_HEADER=1);"
```



## Demo for S3 stage

### 2. Create a S3 stage, Load data from S3 stage to city table.

- b. Create a bucket on S3 called “snowflake-stage-bucket-<your name>”
- c. Create a IAM policy with [this script](#) template, you need to change the bucket name. Give the policy name “snowflake-stage-bucket-policy” or other name you want.
- d. Create an IAM Role 'snowflake-stage-bucket-role' from 'AWS account' for 'This account':
  - i. Require eternal ID <0000>;
  - ii. attach policy 'snowflake-stage-bucket-policy';
  - iii. copy the arn;
- e. create a S3 INTEGRATION, with the ROLE ARN with [this query](#) template on snowflake in Snowflake.
- f. run “*DESC INTEGRATION<your integration name>*” to get **'STORAGE\_AWS\_IAM\_USER\_ARN'** and **'STORAGE\_AWS\_EXTERNAL\_ID'**.
- g. Past the 2 codes to AWS IAM Role, by going to Role -->'trust relationships' , and replace the JSON with [this template](#).





# Demo

## Loading data

- j. Go to snow flake create a FORMAT called CSV\_COMMA with [this query](#). This step is to tell Snowflake, we are going to use the format csv to load data.
- k. grant stage and integration to schema with the following query:
  - *GRANT CREATE STAGE ON SCHEMA ENTERPRISE to ROLE accountadmin;*
  - *GRANT USAGE ON INTEGRATION S3\_INT\_WCD\_LECT1 to ROLE accountadmin;*
- j. Create the STAGE with this query:  
*CREATE OR REPLACE STAGE WCD\_LECT1\_STAGE  
STORAGE\_INTEGRATION = S3\_INT\_WCD\_LECT1  
URL='s3://your bucket name'  
FILE\_FORMAT = CSV\_COMMA;*
- k. Upload the [city.csv](#) file to the bucket;
- l. Check if the file has been in stage *WCD\_LECT1\_STAGE* with query:  
*List @WCD\_LECT1\_STAGE;*
- m. Copy file to city table from stage @WCD\_LECT1\_STAGE with query:  
*COPY INTO walmart\_dev.enterprise.city from @WCD\_LECT1\_STAGE/city.csv  
FILE\_FORMAT =CSV\_COMMA;*

Business Scenarios

Components

Know Snowflake

SnowSQL CLI

**3rd Party Connections**

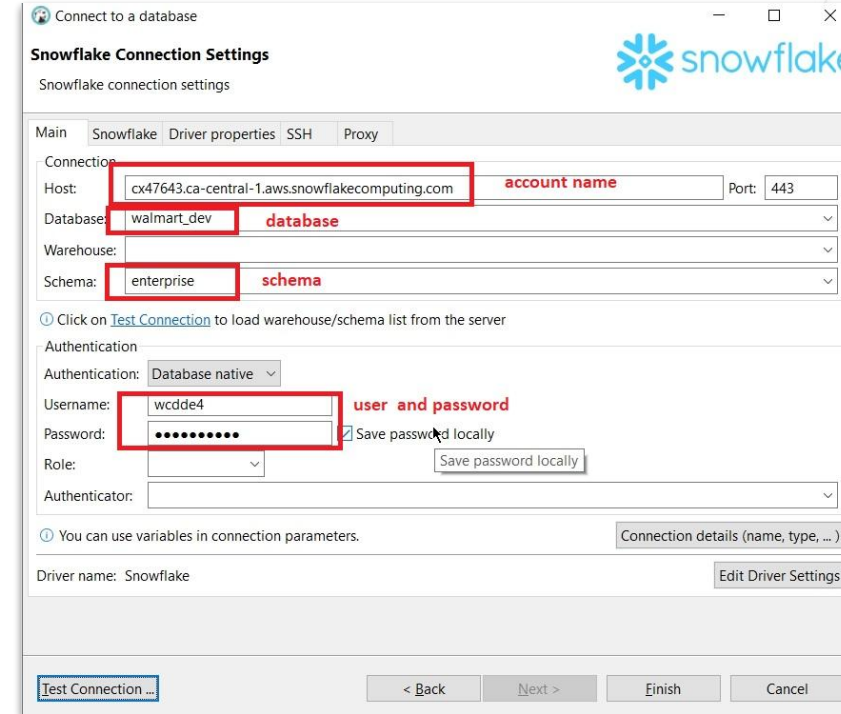
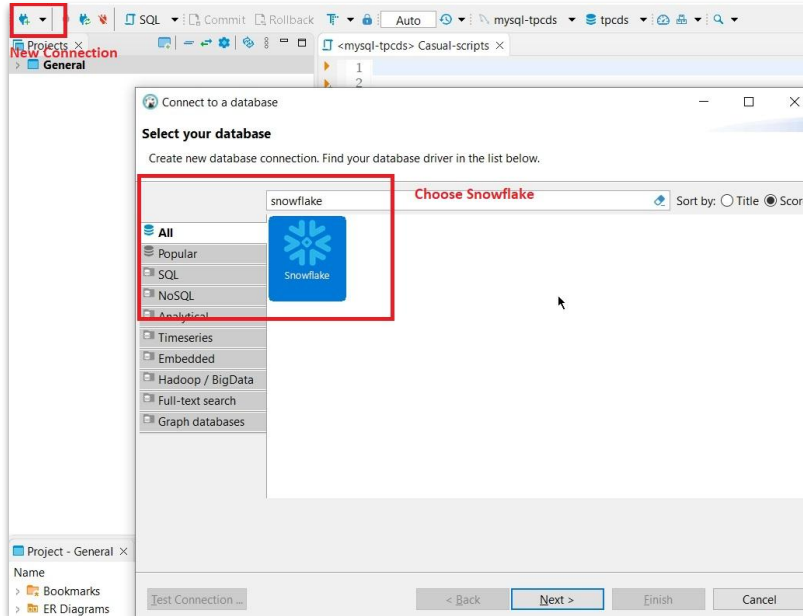
**Agenda.**



# 3rd Party tools Connection

## Loading data

### 1. DBeaver

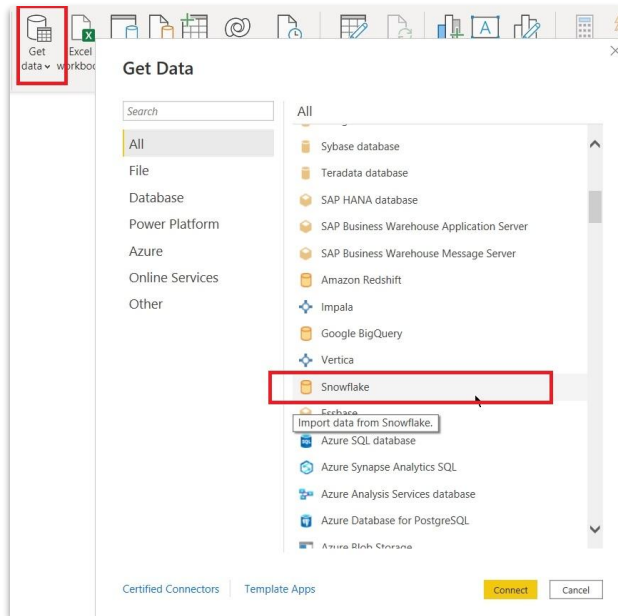




# 3rd Party tools Connection

## Loading data

### 2. PowerBI



**Snowflake**

Server  
cx47643.ca-central-1.aws.snowflakecomputing.com **url link**

Warehouse  
COMPUTE\_WH **warehouse name**

Advanced options  
Specify a text value to use as Role name (optional)  
Example: abc  
Command timeout in seconds (optional)  
Example: 123  
Connection timeout in seconds (optional)  
Example: 123  
Include relationship columns (optional)  
Example: true  
Database (optional)  
walmart\_dev  
SQL statement (optional, requires database)  
Example: sele...

OK Cancel

**Snowflake**

cx47643.ca-central-1.aws.snowflakecomputing.com;...

User name  
WCDE4 **user and password**

Password  
.....

Back Connect Cancel

**Navigator**

Display Options

cx47643.ca-central-1.aws...  
SNOWFLAKE  
SNOWFLAKE\_SAMPLE...  
WALMART\_DEV [2]  
ENTERPRISE [1]  
CITY  
PUBLIC

**choose database + schema + table**

CTY_ID	CTY_NAME
100001	Algiers
100002	Bujumbura
100003	Bangui
100004	Cairo
100005	Addis Ababa
100006	Libreville
100007	Banjul
100008	Conakry
100009	Abidjan
100010	Nairobi
100011	Rabat



# 3rd Party tools Connection

Loading data

## 3. Python

*pip install snowflake-connector-python*

```
import snowflake.connector as sf
import pandas as pd

# make changes as per your credentials
user='snowflake050701'
password = 'Code123456'
account='ozb44782.us-east-1'
database='walmart_dev'
warehouse='COMPUTE_WH'
schema='enterprise'
role='accountadmin'

conn = sf.connect(user = user, password = password, \
                  account = account, warehouse=warehouse, \
                  database=database, schema=schema, role=role)

def run_query(connection,query):
    cursor = conn.cursor()
    cursor.execute(query)
    cursor.close()

# sql = 'show warehouses;'
# run_query(conn, sql)

sql = 'select * from city;'
df = pd.read_sql(sql, conn)
df.head()
```